

Final Year Project Term1 Midterm Report:

1. Review of papers and some other related knowledge

1.1. Research Background and Motivation

1.1.1. Problem Context

Contemporary visual understanding systems excel under closed-set assumptions, where category vocabularies are fixed and fully supervised. However, practical applications increasingly demand open-vocabulary and long-tail recognition, where models must generalize to previously unseen concepts and adapt to evolving taxonomies without exhaustive relabeling. Simultaneously, edge and near-real-time deployment scenarios impose stringent constraints on latency, memory, and energy consumption, creating a dual requirement: broad semantic coverage with efficient inference.

1.1.2. Opportunity from Vision–Language Pretraining

Vision–language models such as CLIP demonstrate that contrastive pretraining on large-scale image–text pairs can produce a shared embedding space where images and natural language descriptions align semantically. This enables zero-shot classification, text-to-image retrieval, and phrase grounding by replacing closed-set classifiers with cosine similarity against textual prompts. The success of prompt ensembling and temperature scaling further indicates that language can serve as a flexible, compositional index into visual semantics.

1.1.3. Gaps in Current Practice

Despite their semantic generality, state-of-the-art vision–language models are computationally heavy, with significant memory footprints and latency that hinder edge deployment. Conversely, lightweight detectors (e.g., YOLO families) meet real-time constraints but remain closed-set and lack robust language grounding. Bridging this gap requires transferring or distilling the semantic alignment learned by large vision–language models into efficient, task-oriented architectures without sacrificing throughput.

1.1.4. Project Positioning

This project targets an efficient open-vocabulary pipeline by: (i) leveraging a frozen text encoder from a pretrained vision–language model to provide a stable semantic reference; (ii) employing a lightweight visual backbone to extract multi-scale image features under strict compute budgets; and (iii) introducing distillation and text-conditioned modulation to align visual features with the text embedding space. The resulting model aims to support zero-shot classification/retrieval and, subsequently, open-vocabulary object detection with practical inference speed.

1.2. Fundamental Theory, Key Techniques, and Representative Related Work

1.2.1. Cross-Modal Alignment via Contrastive Learning

Core idea. Vision–language pretraining aligns images and texts in a shared embedding space using contrastive objectives. Given batches of

image–text pairs, the model maximizes similarity for matched pairs while minimizing it for mismatches, typically with an InfoNCE/CLIP loss, L2-normalized features, and a learnable temperature to control logit scaling.

Encoders. Visual encoders are commonly ResNets or Vision Transformers (ViTs); text encoders are Transformer-based. Pretraining scale and data diversity directly affect coverage of long-tail concepts and compositional generalization.

Prompting and calibration. Prompt ensembling (multiple templates per Representative work. CLIP (Radford et al., 2021) established the modern baseline for zero-shot classification and retrieval. ALIGN (Jia et al., 2021) scaled noisy web data. LiT (Zhai et al., 2022) “locks” the image encoder and tunes the text side to improve transfer.

1.2.2. YOLO Family: Fundamentals, Techniques, and Evolution

Overview and principle. YOLO (You Only Look Once) formulates object detection as direct dense prediction from feature maps, removing the region proposal stage to achieve real-time performance. A single forward pass outputs bounding boxes and class confidences at multiple scales.

Canonical architecture:

Backbone: Efficient feature extraction (e.g., CSPDarknet, ELAN/RepVGG-style blocks, or lightweight MobileNet/EfficientNet variants in custom deployments).

Neck: Multi-scale feature aggregation via FPN/PAN to enhance small/large object localization.

Head: Decoupled classification and box regression. The family evolved from anchor-based (YOLOv3–v5) to anchor-free heads (YOLOv8), simplifying optimization and improving stability.

1.2.3 Knowledge Distillation and Semantic Alignment Strategies

Objectives. Transfer the semantic richness of large VLMs into compact, fast models while preserving alignment with the text space. Distillation types.

Representation regression: Match student visual embeddings to teacher image embeddings via L2 or cosine regression with feature normalization.

Contrastive distillation: Place the student in a contrastive setup using the teacher as a target or teacher-generated pseudo-pairs; improves discriminability.

Multi-level/multi-scale alignment: Align features at several depths and scales to stabilize training and benefit detection heads. Text conditioning.

1.2.4 Method Comparison and Key Differences

We compare four paradigms along supervision, architecture,

efficiency, and deployability: (1) CLIP-like zero-shot classifiers, (2) region-level grounded pretraining (e.g., GLIP/RegionCLIP), (3) open-vocabulary ViT detectors (e.g., OWL-ViT), and (4) lightweight YOLO-based OVD with text alignment (ours). CLIP-style models offer broad semantic coverage and strong zero-shot classification/retrieval but lack explicit region localization and are compute-heavy at high resolutions. Grounded pretraining aligns region–text pairs at scale, yielding robust localization and transfer, yet the training pipelines are complex and teacher-dependent. Open-vocabulary ViT detectors provide clean end-to-end text–region similarity heads and strong accuracy on unseen categories, but their transformer backbones incur higher latency and memory, limiting edge deployment. In contrast, our lightweight YOLO-based design replaces fixed softmax heads with normalized feature projections and cosine similarity to cached text embeddings, adds multi-scale alignment/distillation for region-level semantics, and uses FiLM-based conditioning for query-aware inference. This yields a favourable speed–accuracy trade-off: competitive open-vocabulary detection with substantially lower latency, smaller memory footprint, and streamlined export (ONNX/TensorRT), while preserving zero-shot flexibility.

2. Specifications: Environment Setup and Usage

2.1. Preparation for the Model's Coding

2.1.1. Basic programming environments

Linux, Windows, MacOS all should support our programming.

C++/C, Python, Java and other programming language packages should be downloaded in advance.

2.1.2. Pytorch and other required packages

First, according to our plan, users should install PyTorch and torchvision, as well as small additional dependencies. Plus, it needs a CUDA GPU machine. A dependency list(requirements.txt) will be created when coding. The following then will do the trick: pip install -r requirements.txt

2.2. Usage of the Model

2.2.1. API

In final report, the model module should provide some APIs like available_models() for listing all available ViT and other models, load() for loading model checkpoint files, encode_image() and encode_text() for returning the image/text features encoded by the model, model() for returning the cosine similarity between text and image features, and other useful methods.

2.2.2. Datasets and Demonstrations

In the later report, it will offer a list that shows the model's performance (accuracy, recall, and reference speed) over a wide range of datasets, corresponding datasets and instructions for verifications and demonstrations.

2.2.3. Tasks


This model will be able to do some cross-modal tasks at a low cost, including text-image retrieval, zero-shot image classification, and even open-vocabulary object detection.

3. Model Architecture and Programming Implementations

3.1. A simple model architecture diagram

We propose a dual-stream, lightweight vision-language model that

leverages a frozen CLIP model as a semantic teacher and a lightweight YOLO-based backbone as the student. The model is designed to be end-to-end trainable and supports multiple open-vocabulary tasks without requiring the full CLIP model during inference.



3.2. Architecture:

3.2.1. Vision Backbone

The input image is processed not by the heavy CLIP image encoder, but by a lightweight vision model. A prime candidate is YOLO (e.g., YOLOv5s, YOLOv8n) due to its excellent speed-accuracy trade-off and inherent capability for object-level feature extraction. We will use its backbone and neck for feature extraction, but replace its detection head with a custom projection module.

3.2.2. Text Backbone

The input text prompt (e.g., "a photo of a dog") is processed by the frozen CLIP text encoder (a Vision Transformer, ViT, is typically used in CLIP's *image* encoder; the text encoder is a transformer-based model). Its output is a semantic text embedding.

The detection head is replaced with a custom projection module (e.g., a 2-layer MLP) to map image features into the same embedding space as CLIP.

3.2.3. Fusion Module

Implements Feature-wise Linear Modulation (FiLM) or a cross-attention mechanism to fuse text embeddings with image features.

Text embeddings are used to modulate the intermediate feature

maps of the YOLO backbone, enabling text-guided visual processing. Therefore, the fusion module will consider both distillation loss, and text-image feature alignment. Plus, it will calculate errors between predictions and data for regularization and model fine-tuning. Further research and experiments will be conducted.

3.3. Loss Functions

Distillation Loss: Cosine Embedding Loss or L2 Loss between the projected image embeddings (from YOLO) and the image embeddings from the frozen CLIP image encoder.

Cross-Modal Alignment Loss: Contrastive loss (e.g., InfoNCE) to ensure image and text embeddings are semantically aligned.

3.4. Implementations Strategy

In order to increase work efficiency, it is important to use techniques like de-coupling - reduce interdependence among components.

3.5. Scripts for specific tasks

Several .py files will be provided for specific tasks mentioned before, including cross_modal_retrieval.py, zero_shot_classification.py, open_val_obj_detection.py.

Zero-Shot Classification:

Head: Cosine similarity between image embedding and text class embeddings.

No additional loss; uses similarity scoring.

Text-Image Retrieval:

Same as above, but applied to a database of images and text queries.

Open-Vocabulary Object Detection (OVOD):

Head: Lightweight detection head that outputs region proposals and region-specific embeddings.

Loss: Combination of localization loss (e.g., GIoU) and embedding alignment loss (cosine similarity between region embeddings and text embeddings of class names).

References

- [1] Radford, A., et al. Learning Transferable Visual Models From Natural Language Supervision (CLIP). ICML, 2021.
- [2] Jia, C., et al. Scaling Up Visual and Vision-Language Representation Learning With Noisy Text Supervision (ALIGN). ICML, 2021.
- [3] Zhai, X., et al. LiT: Zero-Shot Transfer with Locked-image Tuning. CVPR, 2022.
- [4] Redmon, J., Farhadi, A. YOLOv3: An Incremental Improvement. arXiv:1804.02767, 2018.
- [5] Bochkovskiy, A., Wang, C.-Y., Liao, H.-Y.M. YOLOv4: Optimal Speed and Accuracy of

Object Detection. CVPR, 2020.

[6] Hinton, G., Vinyals, O., Dean, J. Distilling the Knowledge in a Neural Network. arXiv:1503.02531, 2015.

[7] Tian, Y., Krishnan, D., Isola, P. Contrastive Representation Distillation. ICLR, 2020.

[8] Li, J., et al. GLIP: Grounded Language-Image Pre-training. CVPR, 2022

[9] Minderer, M., et al. Simple Open-Vocabulary Object Detection with Vision Transformers (OWL-ViT). ECCV, 2022; TPAMI, 2023.